

Read the benchmark, not the headline

"AI BEATS DOCTORS AT DIAGNOSIS"

the same result, read three ways

AI scores 94%. Doctors score 89%.

Which test?

1,000 clean textbook images —
not a real OPD queue

Versus whom?

junior residents,
not senior radiologists

Did data leak?

the test images were online —
the model may have seen them

Three questions turn most AI headlines into something much smaller

The headline, and the three questions under it

'Beats humans at X' usually means: **on this test, under these conditions.**

- 1 Ask what the test actually measured — and who chose it.
- 2 Ask what it was compared against. Untrained humans? Experts?
- 3 Ask whether the test data leaked into training. Often nobody checked.

REMEMBER

Three questions turn most AI headlines into something much smaller and more honest.

