


A CNN does not look at the whole image. It slides a **small window** over it.



- 1 The window is a few pixels wide and looks for one pattern — say a vertical edge.
- 2 It slides across the entire image, marking where that pattern appears.
- 3 Stack many windows and you detect many patterns at once.

REMEMBER THIS

Sliding the same small filter everywhere is why CNNs need so little data per feature.